

# Documentation for Empathetic Sentiment Analysis Using TinyLlama-1.1B

## 1. Approach:

The goal of this project was to build an empathetic sentiment analysis model using the **TinyLlama-1.1B-Chat-v1.0** model. The focus was to analyze empathetic dialogues and their relationship to emotions expressed in conversations. The process involved data preparation, model fine-tuning, and performance evaluation, all structured in a well-defined project folder.

## 2. Data Preparation Techniques:

### 1. Dataset:

- The dataset was sourced from the Kaggle dataset [atharvjairath/empathetic-dialogues-facebook-ai](#), containing **64,636 rows** and **7 features**, which included:
  - situation, emotion, empathetic\_dialogues, labels, and 3 unnamed columns.

### 2. Data Cleaning:

- **Missing values** were handled by replacing them with the mode of the respective column and dropping irrelevant rows.
- **Duplicates** were removed to avoid overfitting.
- **Text Standardization:** Removed unwanted characters (e.g., "customer:") and made the text lowercase to maintain consistency.
- Created a new feature, text\_length, to assess any correlation between the length of situation and the associated emotion.

### 3. Exploratory Data Analysis (EDA):

- **Emotion Distribution:** Visualized using countplots to understand the frequency of each emotion.
- **Text Length Distribution:** Visualized using histograms to determine the consistency of text lengths across emotions.
- **Word Cloud Generation:** Created word clouds for situation and empathetic\_dialogues to analyze commonly used words.
- The data was then split into **training** and **testing** datasets, with cleaned and renamed features for convenience.

## 3. Model Choices:

### 1. TinyLlama-1.1B-Chat-v1.0:

- Chosen for its efficiency, as it is lightweight and easy to fine-tune even with limited GPU resources.
- The model is based on causal language modeling, suitable for dialogue generation and empathy-based sentiment analysis.

### 2. Fine-Tuning with BitsAndBytes and PEFT:

- **BitsAndBytes Configuration:** This approach enables the model to be loaded in **4-bit precision**, reducing memory consumption significantly while maintaining efficiency in fine-tuning.
- **PEFT (Performance-Efficient Fine-Tuning):** Specifically, **LoRA** (Low-Rank Adaptation) was used, updating only certain layers (causal\_lm) while freezing the rest of the model's weights to make training more efficient and scalable.

- The fine-tuning was performed using **fp16 precision** for speed optimization.

### 3. Trainer Setup:

- Utilized the **Supervised Fine-Tuning (SFT) Trainer** for model training, using the prepared dataset.
- Fine-tuning was done with **AdamW optimizer**, **learning rate scheduler**, and a total of **2 epochs** for better generalization.

### 4. Model Folder Structure:

#### 1. src Folder:

- **Common Scripts:** Contains `exception_handler.py`, `logger.py`, and utility functions for logging and handling errors.

#### 2. Components Folder:

- **dataingestion.py:** Converts the pre-processed CSV into a Hugging Face dataset format and splits it into training and evaluation datasets.
- **llm\_trainer.py:** Responsible for loading and fine-tuning the TinyLlama model. Utilizes the **BitsAndBytes configuration** and **PEFT** for efficient training.

#### 3. Pipeline Folder:

- **training\_main.py:** Manages the training pipeline, integrating the trainer and data preprocessing.
- **model\_evaluator.py:** Handles evaluation of the fine-tuned model using the test dataset. Calculates accuracy, precision, recall, F1-score, and generates a confusion matrix. It also includes error analysis, displaying the true and predicted labels for each text.
- **inference.py:** Provides an interface for generating conversational responses from the fine-tuned model.

### 5. Challenges Faced:

- **Handling Imbalanced Data:** Emotions were not uniformly distributed, leading to some classes being underrepresented. To handle this, data augmentation techniques could be implemented in future iterations.
- **Memory Constraints:** Fine-tuning large models with limited GPU resources was challenging. Using 4-bit precision and PEFT (LoRA) helped mitigate these issues.
- **Evaluation Metrics:** The evaluation process was hindered by the high class imbalance, and further optimization of evaluation strategies (e.g., using weighted metrics) could improve results.

### 6. Results:

- The model achieved promising results in predicting the **emotion** based on dialogues, with good precision, recall, and F1-scores.
- Confusion matrices showed clear distinctions between major emotions but highlighted some difficulty in predicting more nuanced emotional expressions.

### 7. Reflection & Further Improvements:

#### 1. Improving Cultural Sensitivity:

- The model could be improved by incorporating a more diverse dataset that covers a broader range of cultural contexts and expressions. Current datasets may lack the nuance needed to handle diverse emotional expressions across different cultures.

## 2. Fine-Tuning Techniques:

- More advanced techniques, such as **prompt-tuning** or using domain-specific pre-trained models, could improve the model's ability to detect nuanced emotions better.

## 3. Enhanced Evaluation Metrics